# The Illustrated Transformer

Jay Alammar
Visualizing machine learning one concept at a time.
@JayAlammar on Twitter. YouTube Channel

## Artificial Intelligence

### Creating the Future

Dong-A University

Division of Computer Engineering &
Artificial Intelligence

# References

Main

- https://jalammar.github.io/illustrated-transformer/

# Illustrated Transformer

## Transformer

- A model that uses **attention** to boost the speed with which these models can be trained.
- The Transformers outperforms the Google Neural Machine Translation model in specific tasks.
- The biggest benefit, however, comes from how The Transformer lends itself to **parallelization**. It is in fact Google Cloud's recommendation to use The Transformer as a reference model to use their Cloud TPU offering.

- A TensorFlow implementation of it is available as a part of the Tensor2Tensor package.
- Harvard's NLP group created a guide annotating the paper with PyTorch implementation.

## A High-Level Look

# Illustrated Transformer

## A High-Level Look


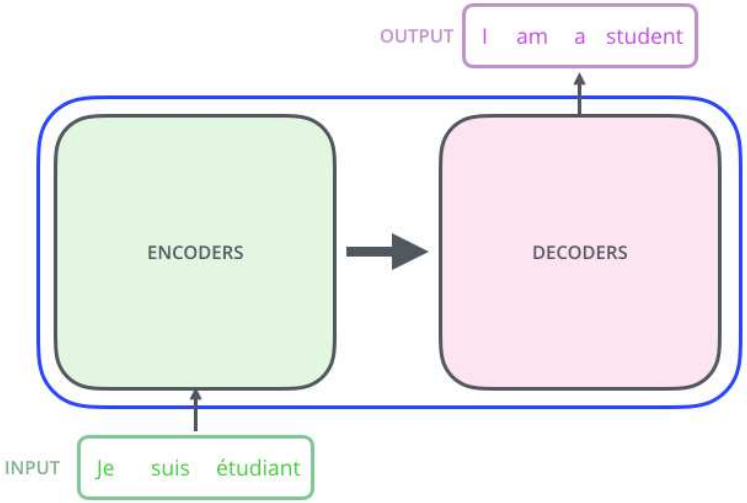

# Illustrated Transformer

## A High-Level Look

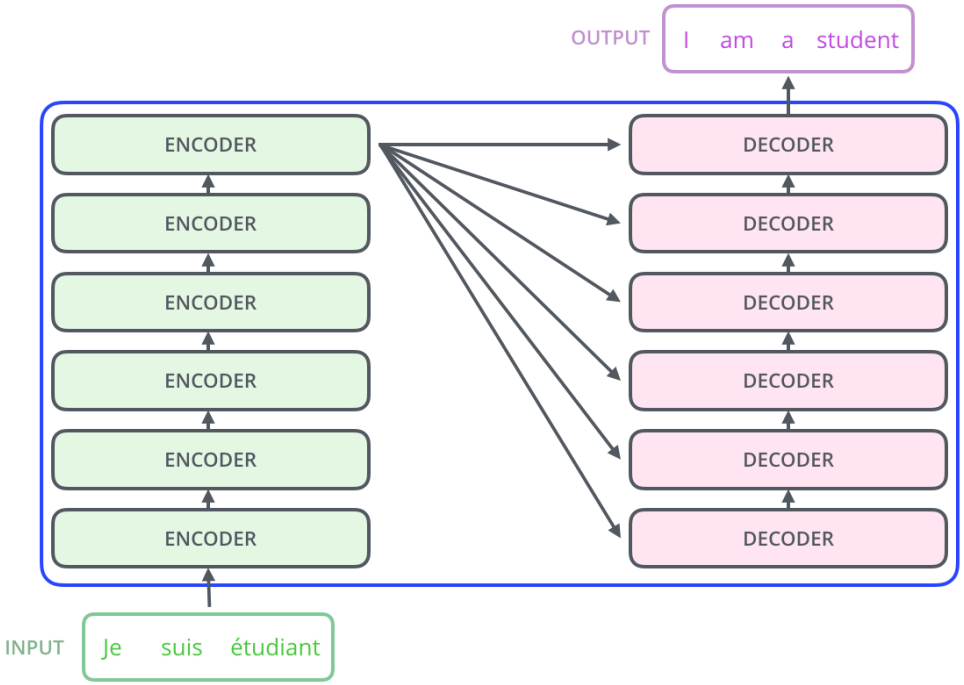


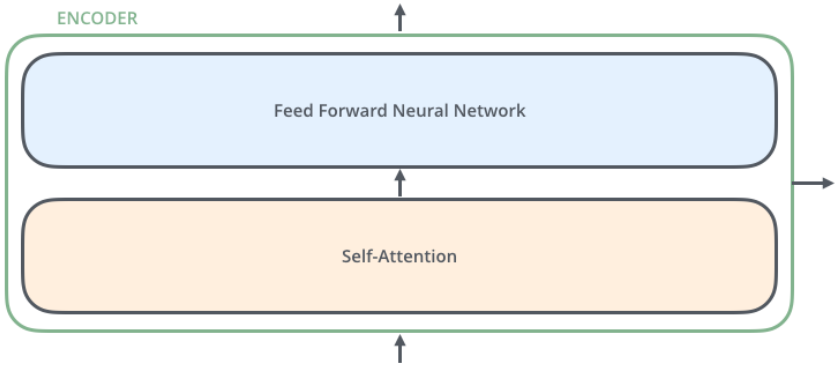


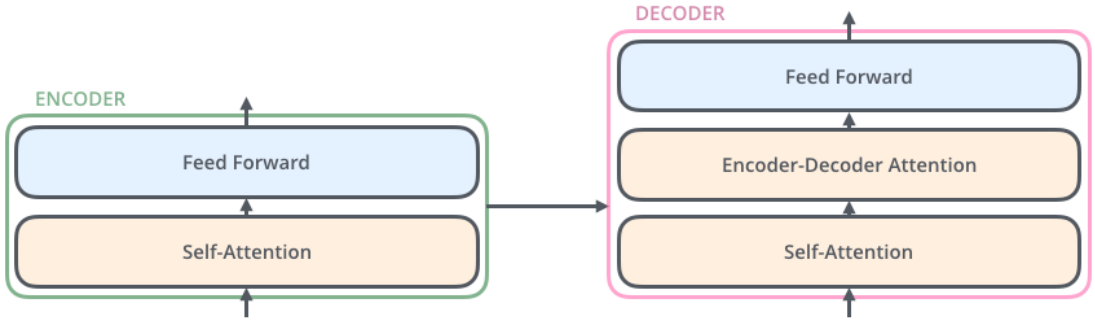

# Illustrated Transformer

## Bringing The Tensors Into The Picture

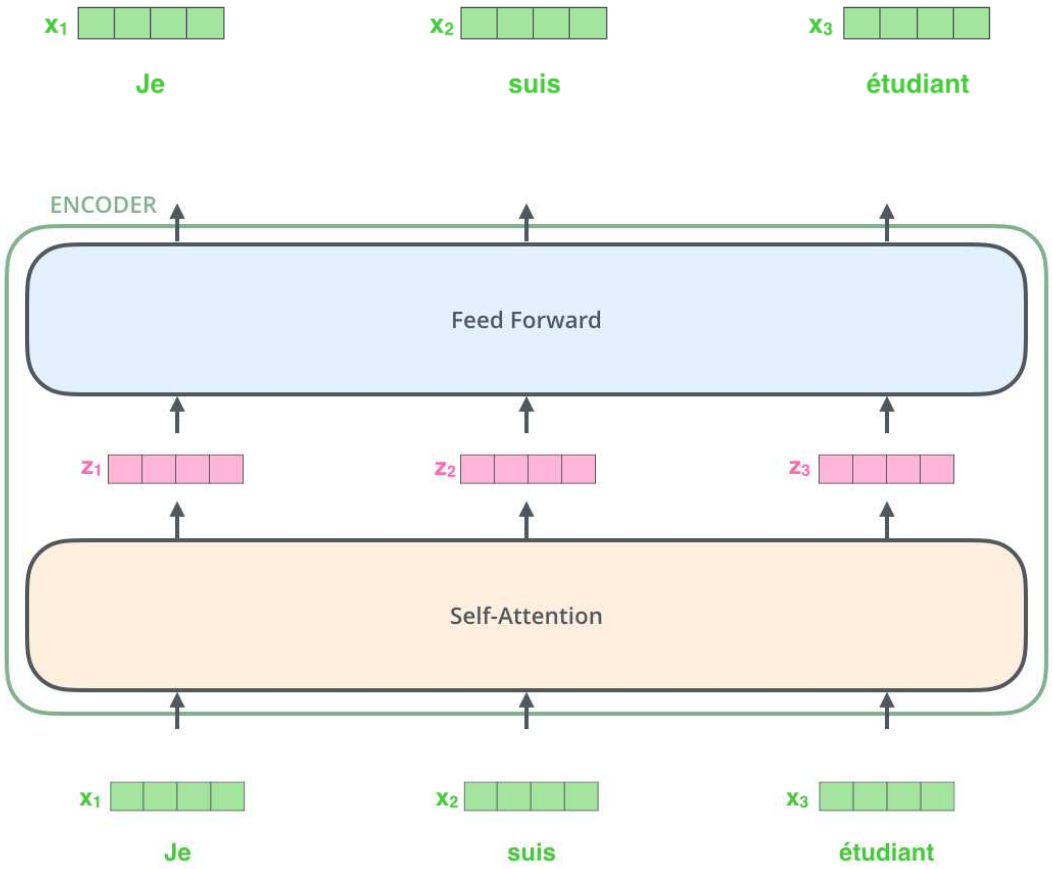


## Bringing The Tensors Into The Picture

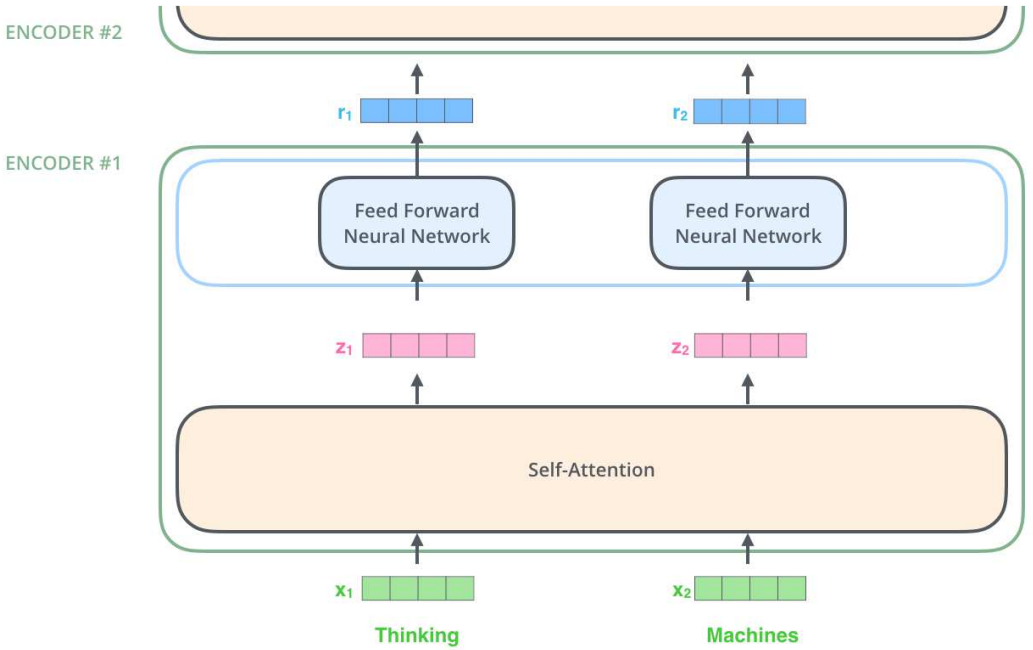

# Illustrated Transformer

## Self-Attention at a High Level

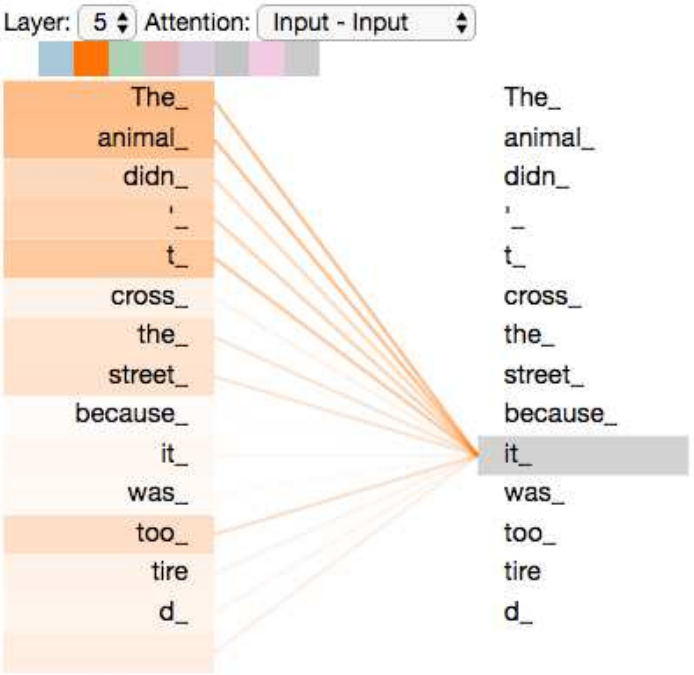


## Self-Attention in Detail

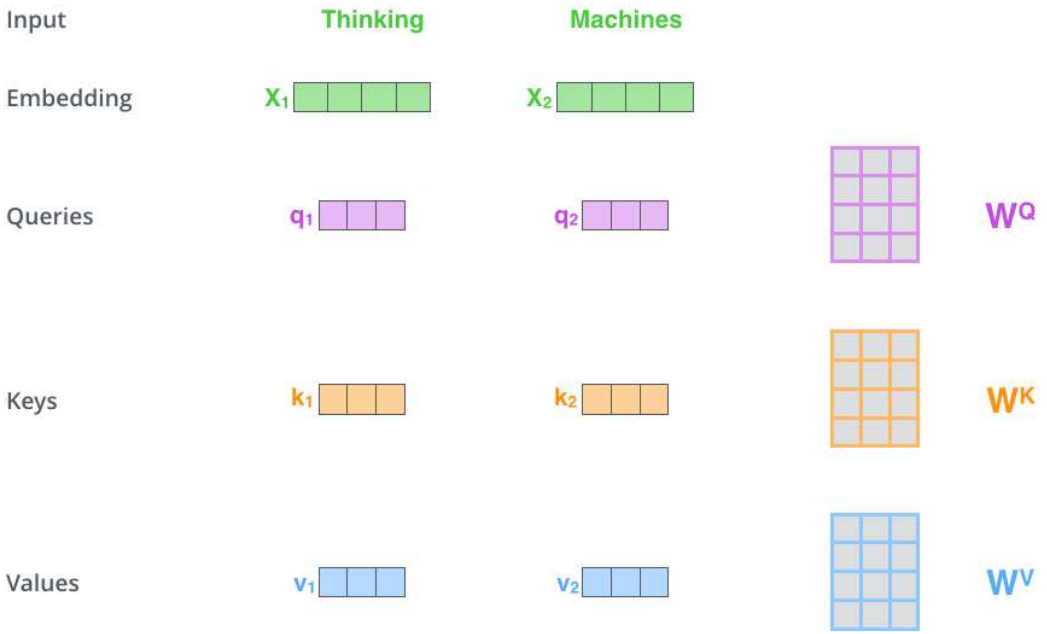

# Illustrated Transformer

## Self-Attention in Detail

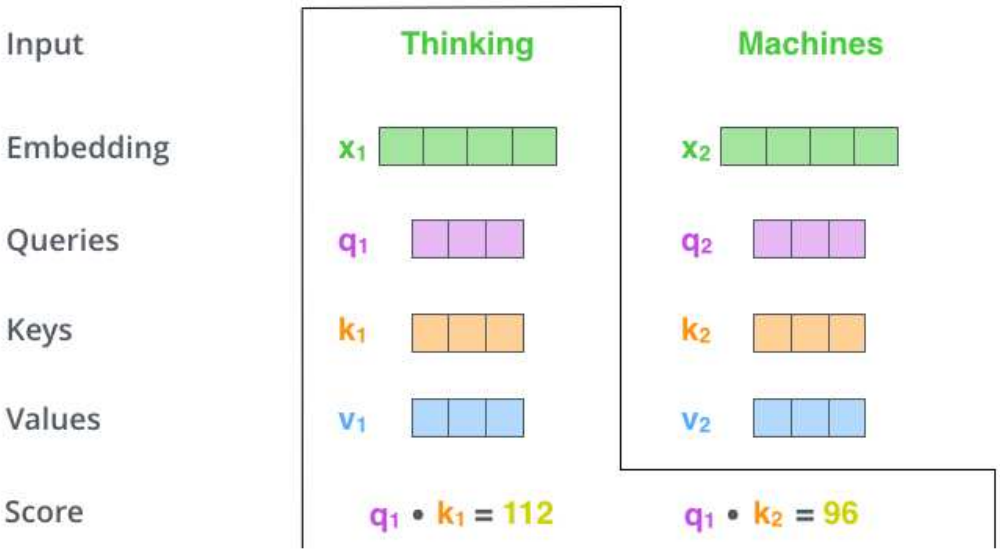


## Self-Attention in Detail

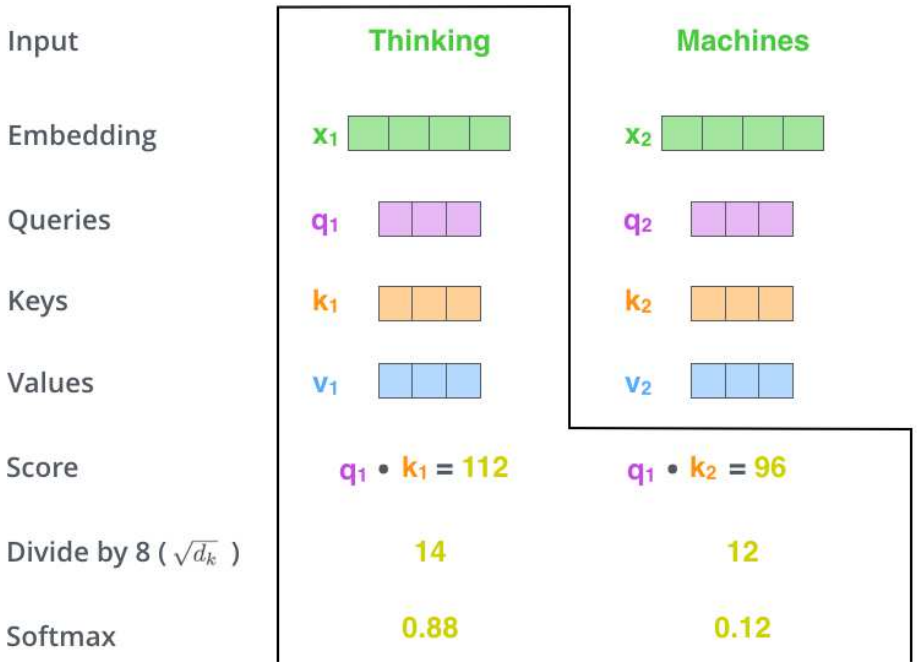

# Illustrated Transformer

## Self-Attention in Detail

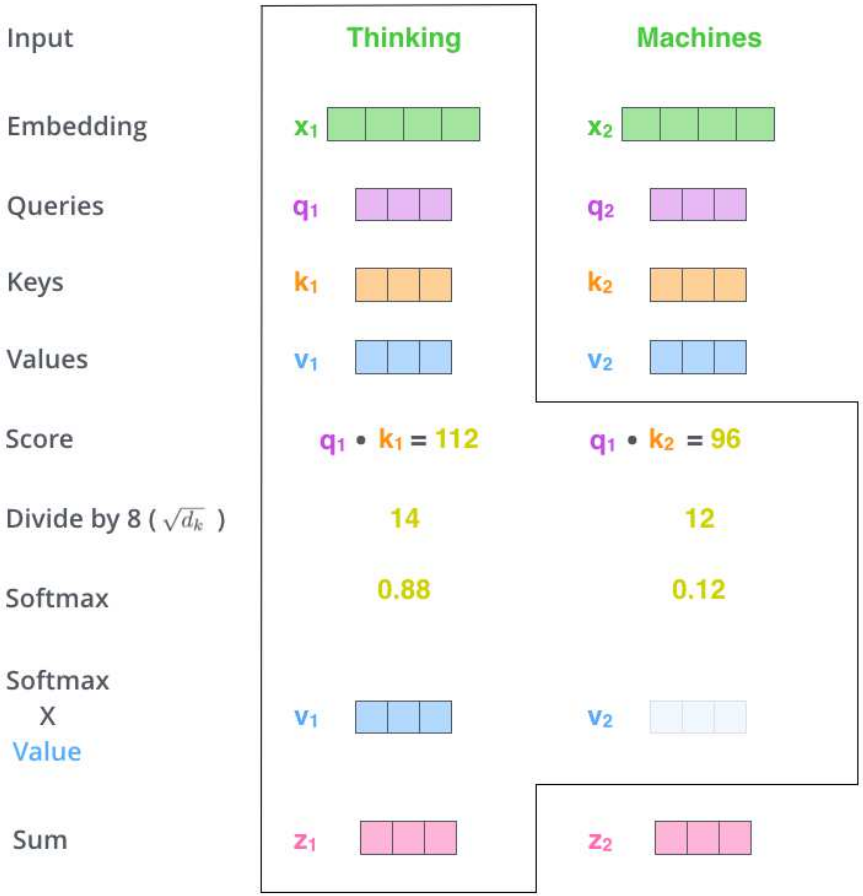


## Matrix Calculation of Self-Attention

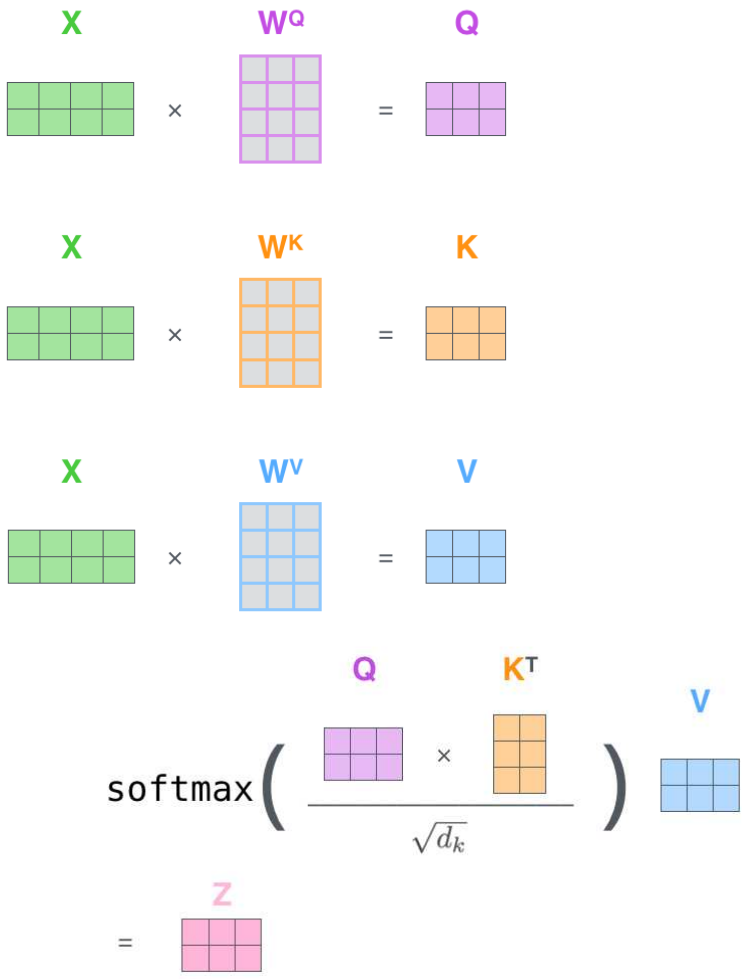

# Illustrated Transformer

## The Beast With Many Heads

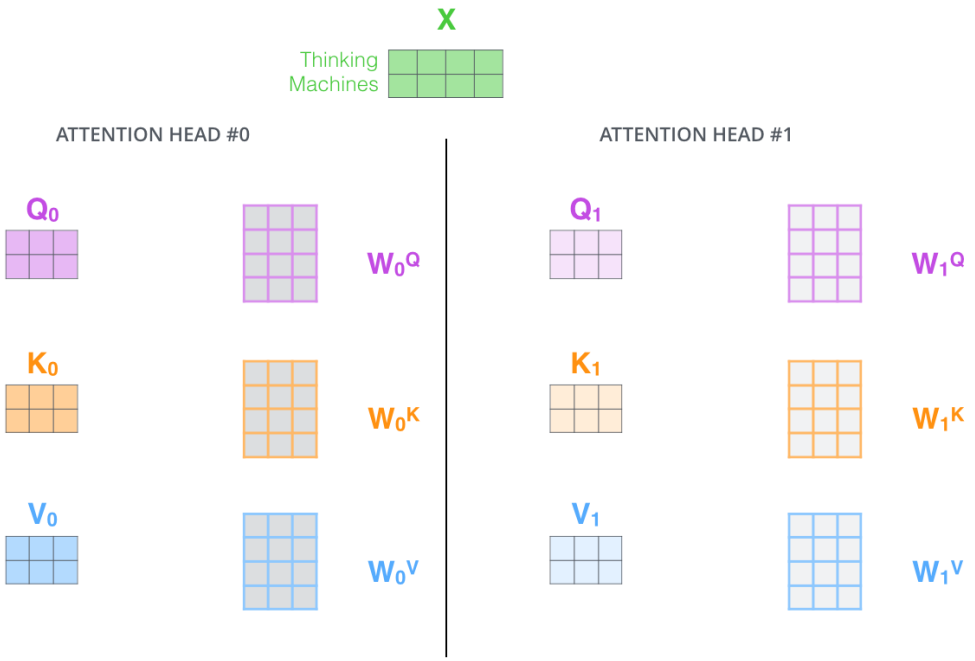


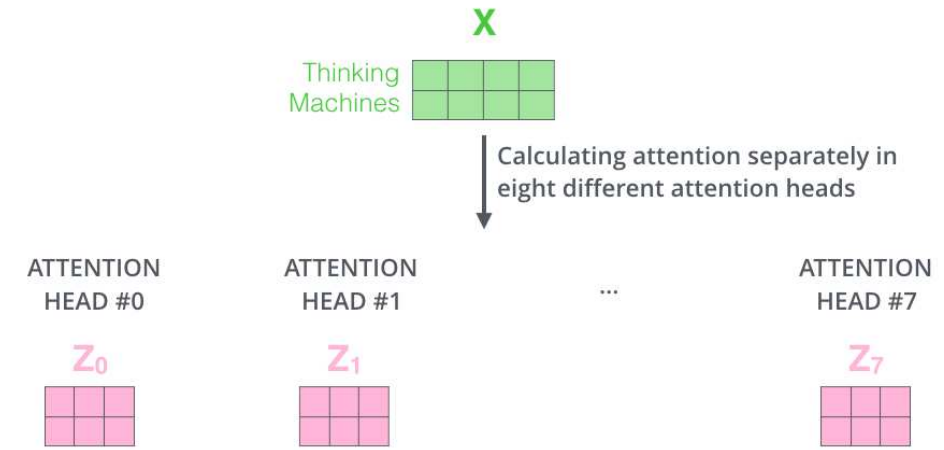

# Illustrated Transformer

## The Beast With Many Heads

1) Concatenate all the attention heads

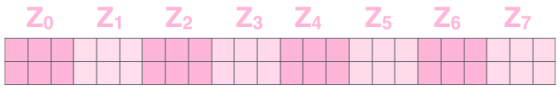


2) Multiply with a weight matrix $W^O$ that was trained jointly with the model

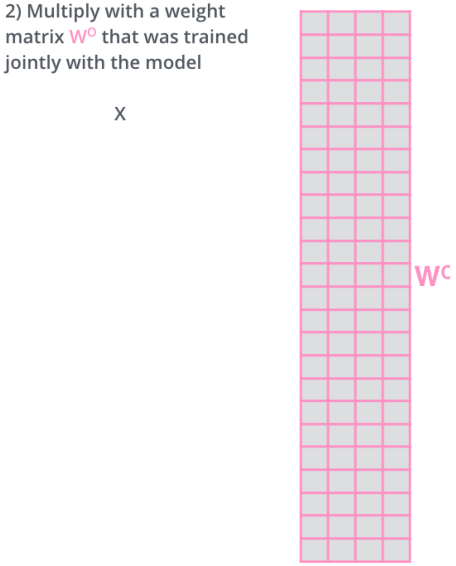


3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN

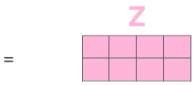


1) This is our input sentence*

Thinking
Machines

2) We embed each word*


* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one


3) Split into 8 heads. We multiply X or R with weight matrices

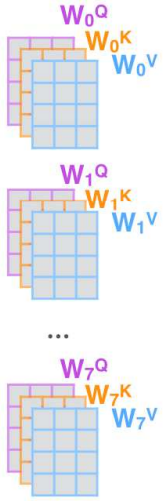


4) Calculate attention using the resulting Q/K/V matrices

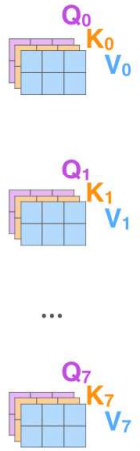


5) Concatenate the resulting Z matrices, then multiply with weight matrix $W^O$ to produce the output of the layer

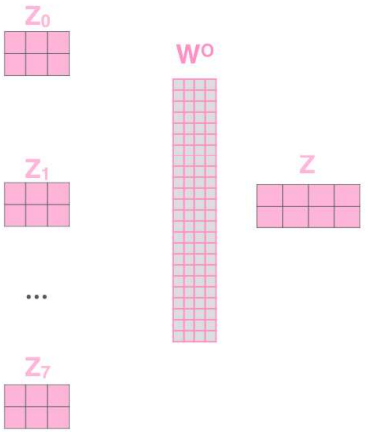

# Illustrated Transformer

## The Beast With Many Heads

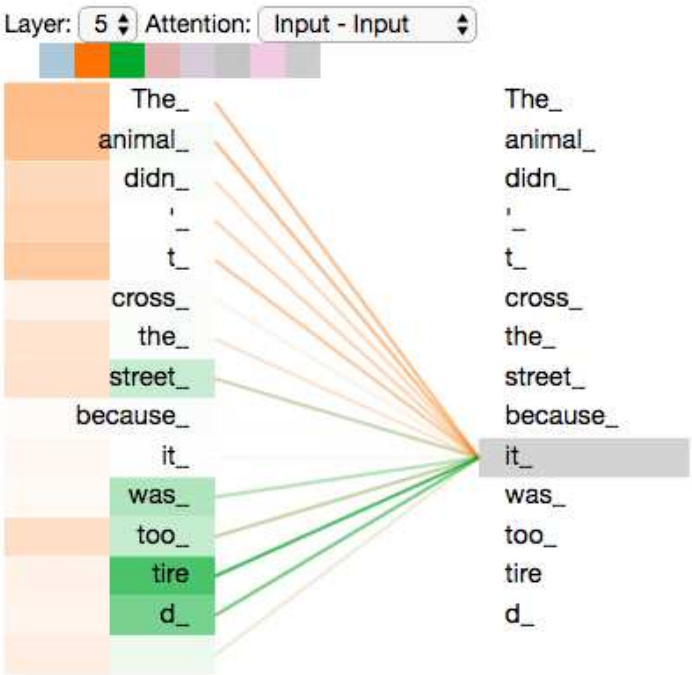


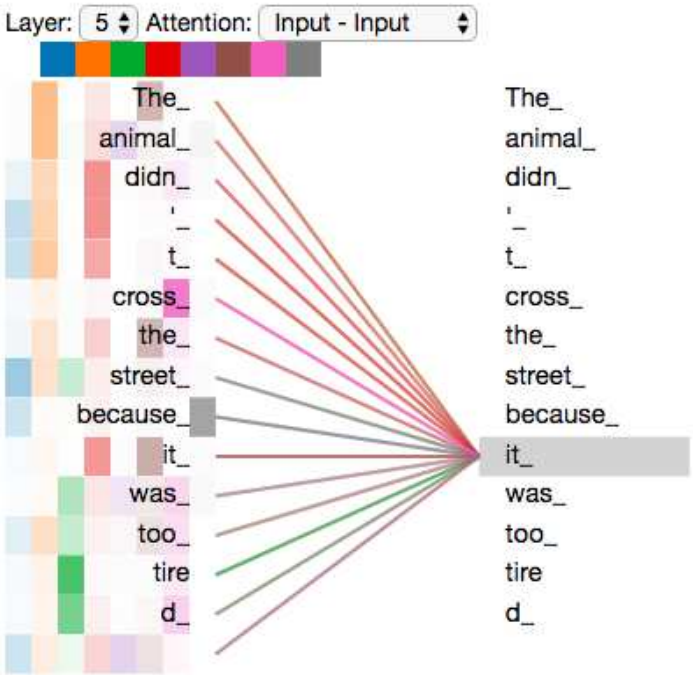

# Illustrated Transformer

## Representing The Order of The Sequence Using Positional Encoding

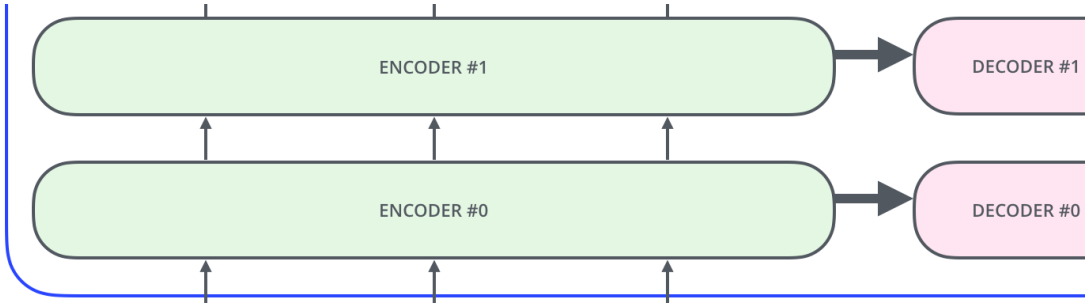


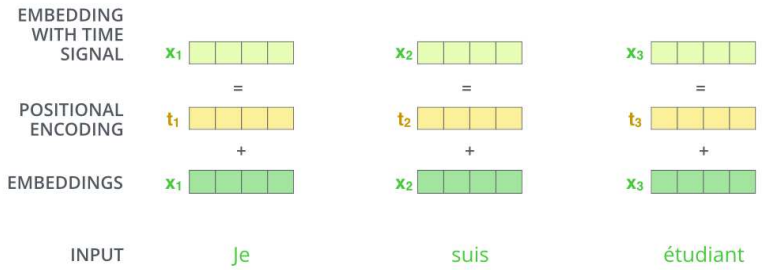


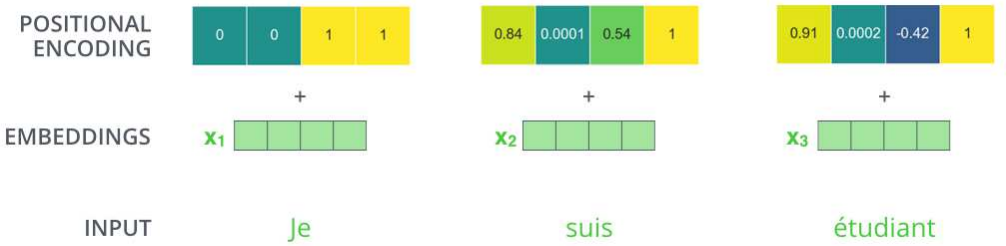

# Illustrated Transformer

## Representing The Order of The Sequence Using Positional Encoding

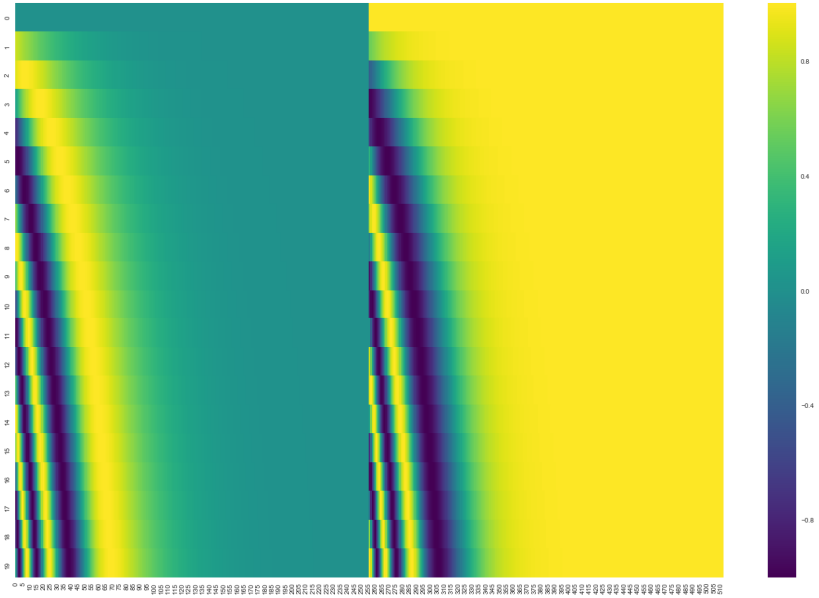

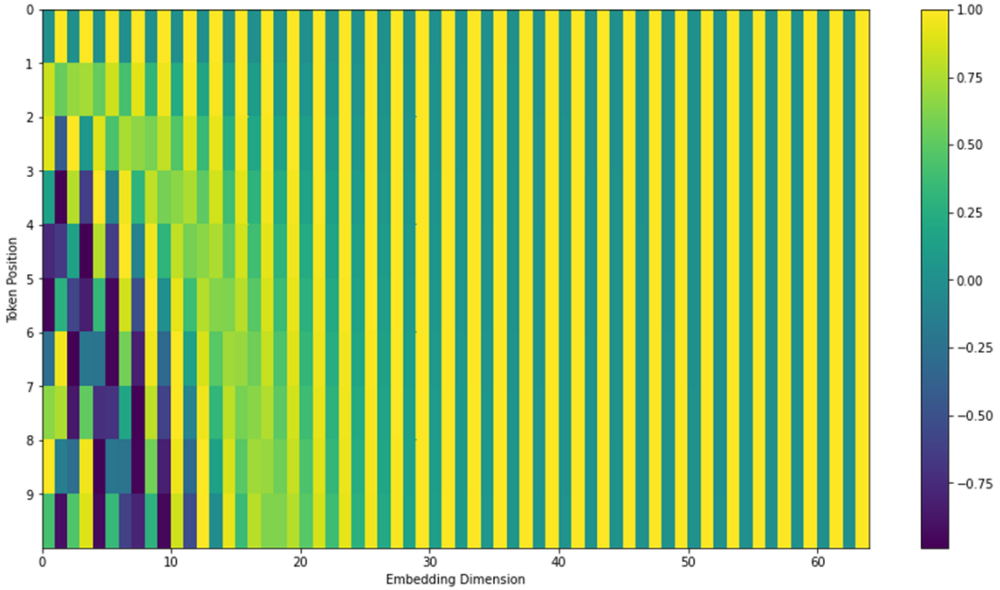

# Illustrated Transformer

## The Residuals

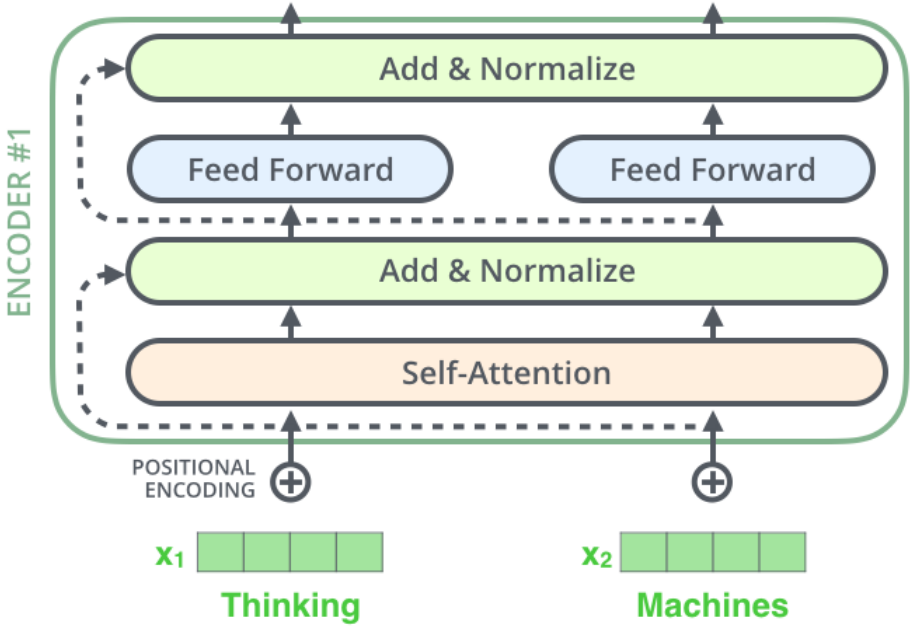


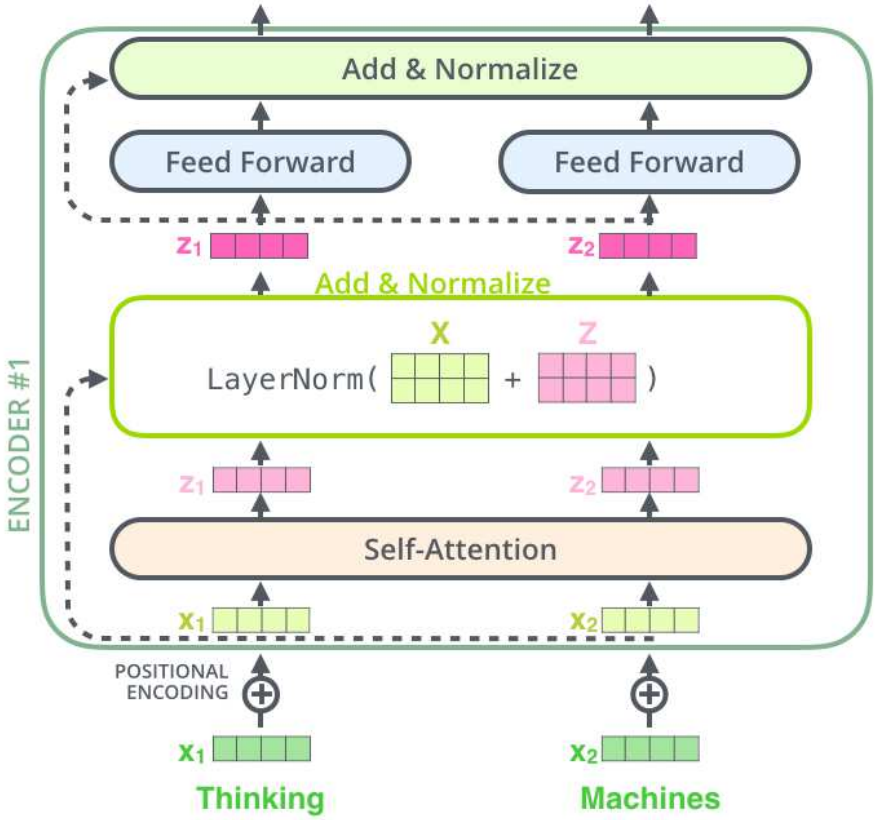

# Illustrated Transformer

## The Residuals

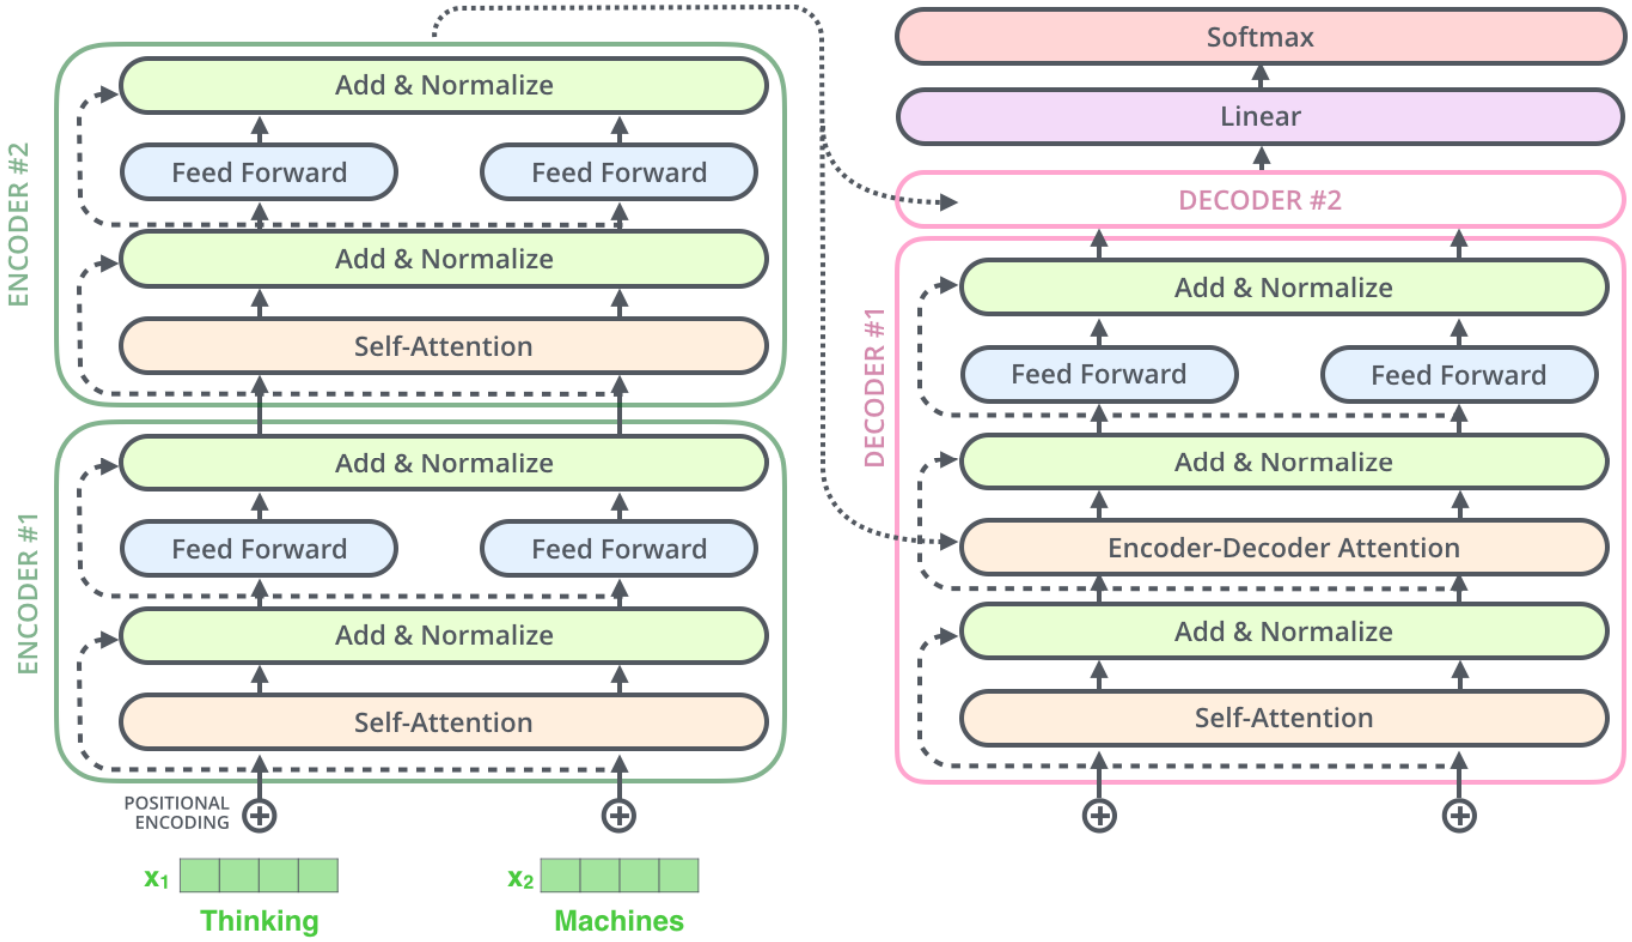

# Illustrated Transformer

## The Decoder Side

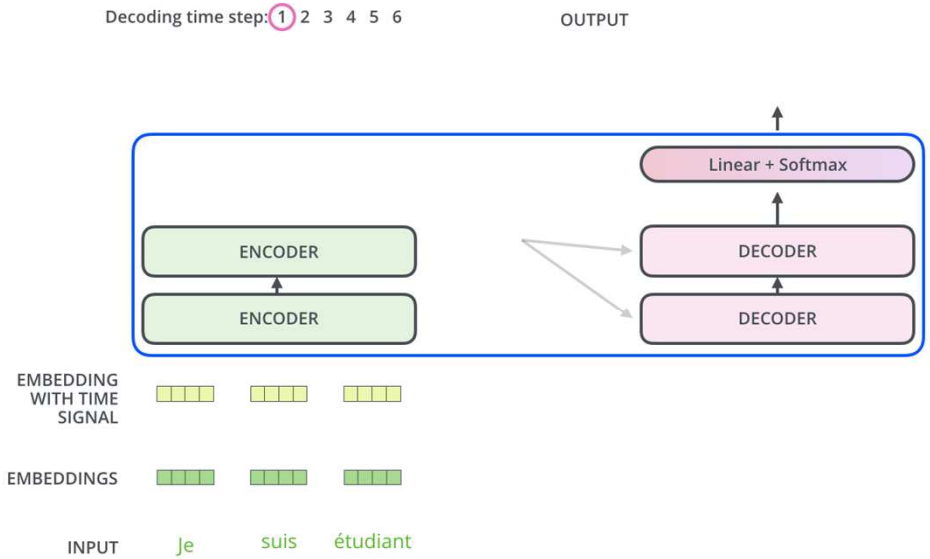


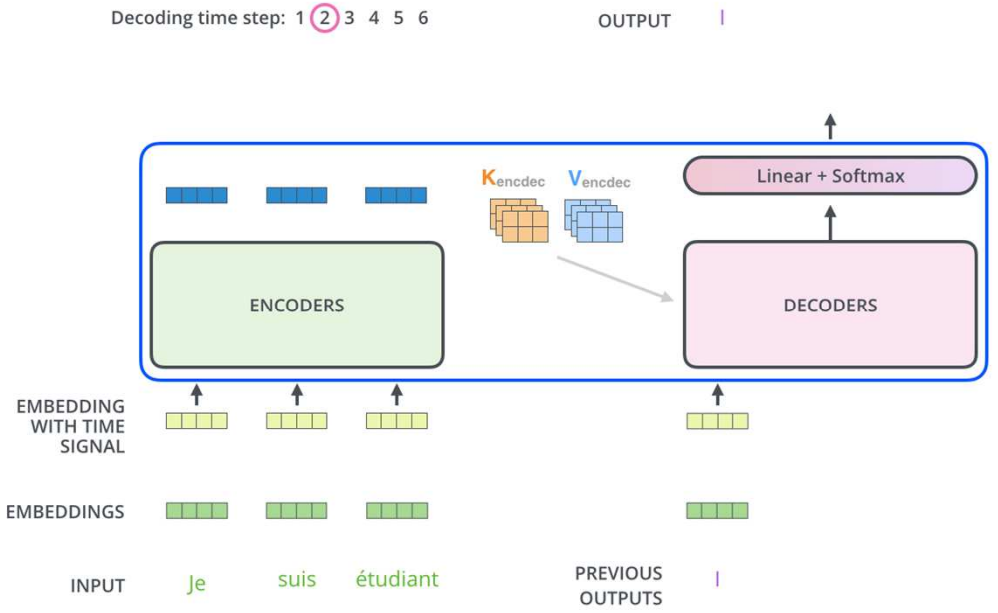

# Illustrated Transformer

## The Final Linear and Softmax Layer

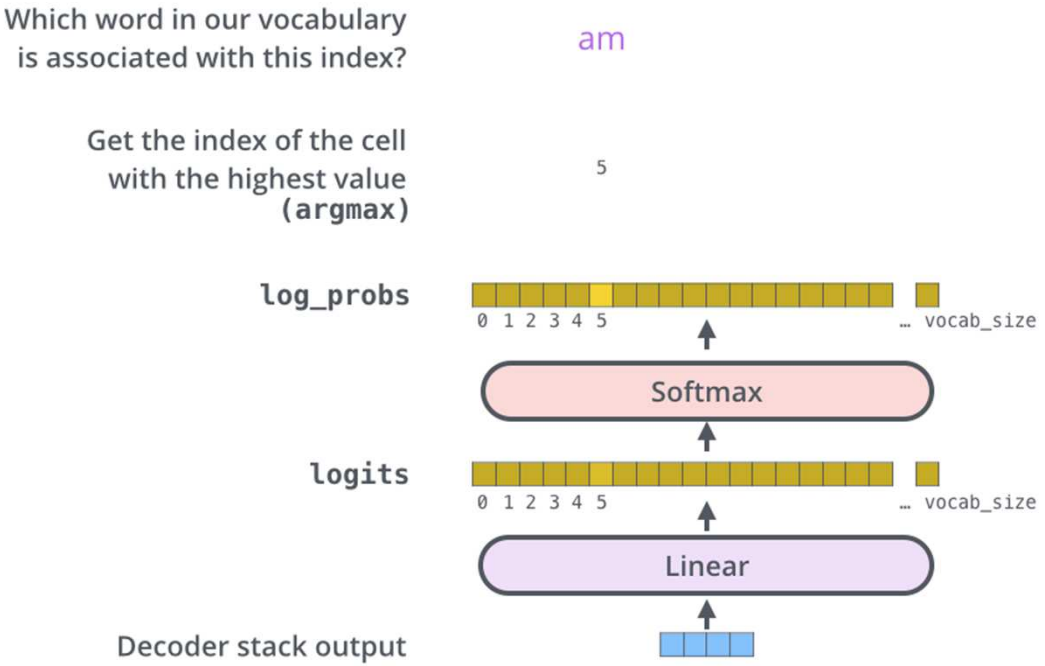


## Recap Of Training

Output Vocabulary

| WORD | a | am | I | thanks | student | <eos> |
|---|---|---|---|---|---|---|
| INDEX | 0 | 1 | 2 | 3 | 4 | 5 |

Output Vocabulary

| WORD | a | am | I | thanks | student | <eos> |
|---|---|---|---|---|---|---|
| INDEX | 0 | 1 | 2 | 3 | 4 | 5 |

One-hot encoding of the word “am”

| 0.0 | 1.0 | 0.0 | 0.0 | 0.0 | 0.0 |
|---|---|---|---|---|---|

# Illustrated Transformer

## The Loss Function